

LetsGrowMore (LGMVIP) - "DATA SCIENCE INTERN"

Name - Shiva Dagdu Mehenge

More Advanced Level Task

TASK-10- : ML Facial recognition to detect mood and suggest songs accordingly

Dataset link : <https://www.kaggle.com/msambare/fer2013>
[\(https://www.kaggle.com/msambare/fer2013\)](https://www.kaggle.com/msambare/fer2013)

```
In [2]: 1 pip install scikit-image
```

```
Collecting scikit-image
  Downloading scikit_image-0.19.2-cp310-cp310-win_amd64.whl (12.6 MB)
    ----- 12.6/12.6 MB 1.9 MB/s eta 0:00:00
Requirement already satisfied: numpy>=1.17.0 in c:\users\sbpat\appdata\local\programs\python\python310\lib\site-packages (from scikit-image) (1.22.0)
Requirement already satisfied: packaging>=20.0 in c:\users\sbpat\appdata\local\programs\python\python310\lib\site-packages (from scikit-image) (21.3)
Requirement already satisfied: scipy>=1.4.1 in c:\users\sbpat\appdata\local\programs\python\python310\lib\site-packages (from scikit-image) (1.7.3)
Collecting PyWavelets>=1.1.1
  Downloading PyWavelets-1.3.0-cp310-cp310-win_amd64.whl (4.2 MB)
    ----- 4.2/4.2 MB 2.8 MB/s eta 0:00:00
Collecting tifffile>=2019.7.26
  Downloading tifffile-2022.4.26-py3-none-any.whl (191 kB)
    ----- 191.8/191.8 KB 3.9 MB/s eta 0:00:00
Requirement already satisfied: pillow!=7.1.0,!>7.1.1,!>8.3.0,>=6.1.0 in c:\users\sbpat\appdata\local\programs\python\python310\lib\site-packages (from scikit-image) (9.0.0)
Collecting networkx>=2.2
  Downloading networkx-2.8-py3-none-any.whl (2.0 MB)
    ----- 2.0/2.0 MB 2.7 MB/s eta 0:00:00
Collecting imageio>=2.4.1
  Downloading imageio-2.18.0-py3-none-any.whl (3.4 MB)
    ----- 3.4/3.4 MB 3.1 MB/s eta 0:00:00
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in c:\users\sbpat\appdata\local\programs\python\python310\lib\site-packages (from packaging>=20.0->scikit-image) (3.0.7)
Installing collected packages: tifffile, PyWavelets, networkx, imageio, scikit-image
Successfully installed PyWavelets-1.3.0 imageio-2.18.0 networkx-2.8 scikit-image-0.19.2 tifffile-2022.4.26
Note: you may need to restart the kernel to use updated packages.
```

Importing Libraries

In [3]:

```
1 #Import Libraries/Packages
2 import numpy as np
3 import cv2
4 import os
5 import random
6 from skimage.io import imread
7 import pandas as pd
8 import seaborn as sns
9 import tensorflow as tf
10 from tensorflow.keras.models import Sequential
11 from tensorflow.keras.layers import Dense, Dropout, Flatten, BatchNormalization, Conv2D, MaxPooling2D, Activation
12 from tensorflow.keras.optimizers import Adam
13 from tensorflow.keras.preprocessing.image import ImageDataGenerator
14 from tensorflow.keras.callbacks import ReduceLROnPlateau, ModelCheckpoint
15 from tensorflow.keras.models import load_model
16 import matplotlib.pyplot as plt
17 from IPython.display import Audio
18
```

In [4]:

```
1 print("TensorFlow version : ", tf.__version__)
2 print("Num GPUs Available: ", len(tf.config.experimental.list_physical_devices('GPU')))
```

TensorFlow version : 2.8.0

Num GPUs Available: 0

Data Visualization

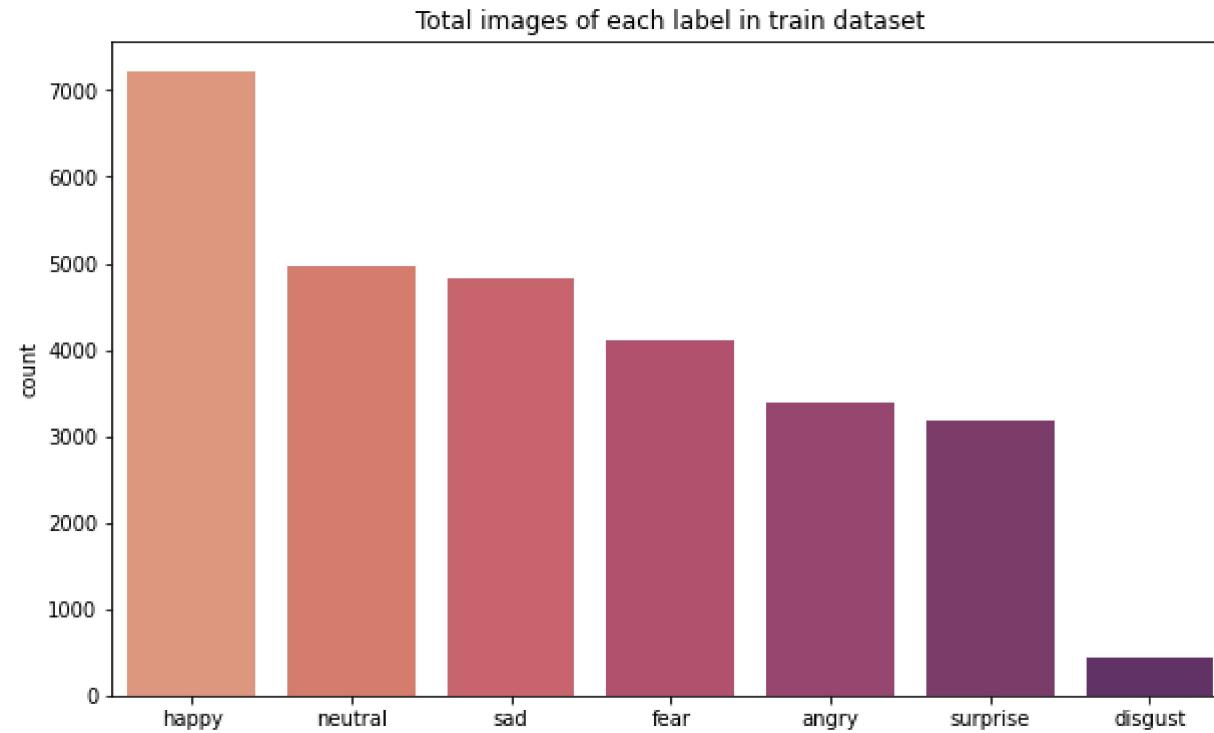
In [13]:

```
1 # Data Visualization
2 from skimage.io import imread
3 train_dir = "train/"
4 test_dir = "test/"
5 total_labels = len(os.listdir(train_dir))
6
7 fig, ax = plt.subplots(nrows=5, ncols=total_labels, figsize=(35, 25))
8 for x in range(5):
9     for y,v in zip(range(total_labels),os.listdir(train_dir)):
10         ax[x][y].imshow(imread(train_dir+v+'/'+os.listdir(train_dir+v)[x]), cmap='gray')
11
12 plt.show()
```



In [14]:

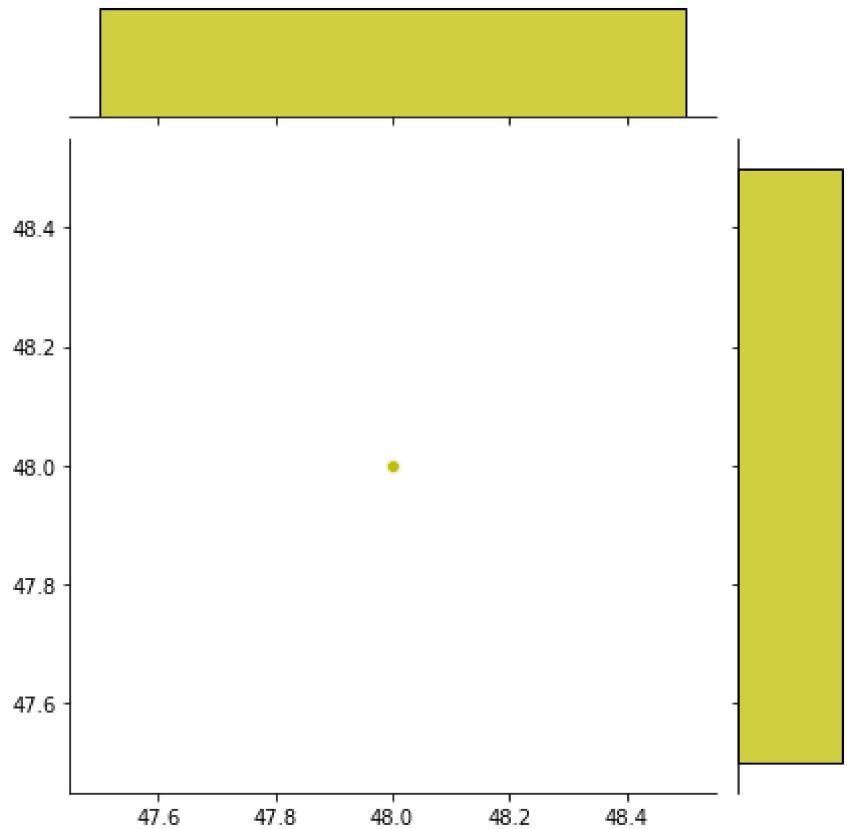
```
1 df = {}
2 for i in os.listdir(train_dir):
3     directory = train_dir + i
4     df[i] = len(os.listdir(directory))
5 df = pd.DataFrame(df, index=["total"]).transpose().sort_values("total", ascending=False)
6
7 plt.figure(figsize=(10,6))
8 sns.barplot(x=df.index, y="total", palette="flare", data=df)
9 plt.ylabel("count")
10 plt.title("Total images of each label in train dataset")
11 plt.show()
12
```



In [15]:

```
1 happy = os.listdir(train_dir+'happy/')
2 dim1, dim2 = [], []
3
4 for img_filename in happy:
5     img = imread(train_dir+'happy/'+img_filename)
6     d1, d2 = img.shape
7     dim1.append(d1)
8     dim2.append(d2)
9
10 img_shape = (int(np.mean(dim1)), int(np.mean(dim2)), 1)
11 sns.jointplot(dim1, dim2,color='y')
12 plt.show()
```

C:\Users\sbpat\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(



In [16]:

```
1 # Data Preprocessing
2 train_gen = ImageDataGenerator(rescale=1/255,
3                                rotation_range=40,
4                                width_shift_range=0.2,
5                                height_shift_range=0.2,
6                                shear_range=0.2,
7                                zoom_range=0.2,
8                                horizontal_flip=True,
9                                fill_mode='nearest')
10
11 test_gen = ImageDataGenerator(rescale=1/255)
12
13 img_shape = (int(np.mean(dim1)), int(np.mean(dim2)), 1)
14
15 train_generator = train_gen.flow_from_directory(directory=train_dir,
16                                                 target_size=(img_shape[0], img_shape[1]),
17                                                 color_mode='grayscale',
18                                                 batch_size=64,
19                                                 class_mode='categorical',
20                                                 shuffle=True)
21
22 test_generator = test_gen.flow_from_directory(directory=test_dir,
23                                                 target_size=(img_shape[0], img_shape[1]),
24                                                 color_mode='grayscale',
25                                                 batch_size=64,
26                                                 class_mode='categorical',
27                                                 shuffle=False)
28
```

Found 28111 images belonging to 7 classes.

Found 7178 images belonging to 7 classes.

Creating the Model

In [17]:

```
1 # Create the Model
2 model = Sequential()
3
4 model.add(Conv2D(filters=64, kernel_size=(3,3), padding='same', activation='relu', input_shape=img_shape)
5 model.add(BatchNormalization())
6 model.add(MaxPooling2D(pool_size=(2,2)))
7 model.add(Dropout(0.2))
8
9 model.add(Conv2D(filters=128, kernel_size=(3,3), padding='same', activation='relu'))
10 model.add(BatchNormalization())
11 model.add(MaxPooling2D(pool_size=(2,2)))
12 model.add(Dropout(0.2))
13
14 model.add(Conv2D(filters=512, kernel_size=(3,3), padding='same', activation='relu'))
15 model.add(BatchNormalization())
16 model.add(MaxPooling2D(pool_size=(2,2)))
17 model.add(Dropout(0.2))
18
19 model.add(Conv2D(filters=512, kernel_size=(3,3), padding='same', activation='relu'))
20 model.add(BatchNormalization())
21 model.add(MaxPooling2D(pool_size=(2,2)))
22 model.add(Dropout(0.2))
23
24 model.add(Flatten())
25
26 model.add(Dense(512, activation='relu'))
27 model.add(Dropout(0.2))
28
29 model.add(Dense(1024, activation='relu'))
30 model.add(Dropout(0.2))
31
32 model.add(Dense(units=len(os.listdir(train_dir)), activation='softmax'))
33
34 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 48, 48, 64)	640
batch_normalization (BatchN ormalization)	(None, 48, 48, 64)	256
max_pooling2d (MaxPooling2D)	(None, 24, 24, 64)	0
dropout (Dropout)	(None, 24, 24, 64)	0
conv2d_1 (Conv2D)	(None, 24, 24, 128)	73856
batch_normalization_1 (Bac hNormalization)	(None, 24, 24, 128)	512
max_pooling2d_1 (MaxPooling 2D)	(None, 12, 12, 128)	0
dropout_1 (Dropout)	(None, 12, 12, 128)	0
conv2d_2 (Conv2D)	(None, 12, 12, 512)	590336
batch_normalization_2 (Bac hNormalization)	(None, 12, 12, 512)	2048
max_pooling2d_2 (MaxPooling 2D)	(None, 6, 6, 512)	0
dropout_2 (Dropout)	(None, 6, 6, 512)	0
conv2d_3 (Conv2D)	(None, 6, 6, 512)	2359808
batch_normalization_3 (Bac hNormalization)	(None, 6, 6, 512)	2048
max_pooling2d_3 (MaxPooling 2D)	(None, 3, 3, 512)	0
dropout_3 (Dropout)	(None, 3, 3, 512)	0

```
flatten (Flatten)           (None, 4608)          0
dense (Dense)               (None, 512)           2359808
dropout_4 (Dropout)         (None, 512)           0
dense_1 (Dense)             (None, 1024)          525312
dropout_5 (Dropout)         (None, 1024)          0
dense_2 (Dense)             (None, 7)              7175
=====
Total params: 5,921,799
Trainable params: 5,919,367
Non-trainable params: 2,432
```

Training the model

```
In [18]: 1 # Train the model
2 model.compile(optimizer=Adam(learning_rate=0.0001,
3                             decay=1e-6),
4                 loss='categorical_crossentropy',
5                 metrics=['accuracy'])
```

In [19]:

```
1 steps_per_epoch = train_generator.n // train_generator.batch_size
2 validation_steps = test_generator.n // test_generator.batch_size
3 num_epochs = 20
4
5 history = model.fit(train_generator,
6                     epochs=num_epochs,
7                     verbose=1,
8                     #callbacks=callbacks,
9                     validation_data=test_generator,
10                    steps_per_epoch=steps_per_epoch,
11                    validation_steps=validation_steps)
```

Epoch 1/20
439/439 [=====] - 1666s 4s/step - loss: 1.8728 - accuracy: 0.2377 - val_loss: 2.444
0 - val_accuracy: 0.2479
Epoch 2/20
439/439 [=====] - 1669s 4s/step - loss: 1.7967 - accuracy: 0.2564 - val_loss: 1.787
5 - val_accuracy: 0.2679
Epoch 3/20
439/439 [=====] - 1725s 4s/step - loss: 1.7828 - accuracy: 0.2593 - val_loss: 1.719
2 - val_accuracy: 0.2916
Epoch 4/20
439/439 [=====] - 1613s 4s/step - loss: 1.7722 - accuracy: 0.2663 - val_loss: 1.771
3 - val_accuracy: 0.2949
Epoch 5/20
439/439 [=====] - 1241s 3s/step - loss: 1.7545 - accuracy: 0.2841 - val_loss: 1.713
9 - val_accuracy: 0.3083
Epoch 6/20
439/439 [=====] - 1166s 3s/step - loss: 1.7321 - accuracy: 0.2978 - val_loss: 1.730
4 - val_accuracy: 0.3225
Epoch 7/20
439/439 [=====] - 1177s 3s/step - loss: 1.7065 - accuracy: 0.3117 - val_loss: 1.590
7 - val_accuracy: 0.3712
Epoch 8/20
439/439 [=====] - 1174s 3s/step - loss: 1.6698 - accuracy: 0.3302 - val_loss: 1.604
5 - val_accuracy: 0.3733
Epoch 9/20
439/439 [=====] - 1173s 3s/step - loss: 1.6355 - accuracy: 0.3503 - val_loss: 1.484
1 - val_accuracy: 0.4248
Epoch 10/20
439/439 [=====] - 1028s 2s/step - loss: 1.5945 - accuracy: 0.3724 - val_loss: 1.452
6 - val_accuracy: 0.4344
Epoch 11/20
439/439 [=====] - 1086s 2s/step - loss: 1.5521 - accuracy: 0.3937 - val_loss: 1.428
6 - val_accuracy: 0.4577
Epoch 12/20
439/439 [=====] - 1036s 2s/step - loss: 1.5221 - accuracy: 0.4070 - val_loss: 1.581
3 - val_accuracy: 0.4134
Epoch 13/20
439/439 [=====] - 984s 2s/step - loss: 1.4914 - accuracy: 0.4213 - val_loss: 1.3892
- val_accuracy: 0.4605
Epoch 14/20
439/439 [=====] - 985s 2s/step - loss: 1.4579 - accuracy: 0.4354 - val_loss: 1.3432
- val_accuracy: 0.4710
Epoch 15/20

```
439/439 [=====] - 983s 2s/step - loss: 1.4303 - accuracy: 0.4454 - val_loss: 1.3578
- val_accuracy: 0.4738
Epoch 16/20
439/439 [=====] - 988s 2s/step - loss: 1.4166 - accuracy: 0.4552 - val_loss: 1.3742
- val_accuracy: 0.4764
Epoch 17/20
439/439 [=====] - 991s 2s/step - loss: 1.3908 - accuracy: 0.4646 - val_loss: 1.2980
- val_accuracy: 0.5015
Epoch 18/20
439/439 [=====] - 989s 2s/step - loss: 1.3833 - accuracy: 0.4686 - val_loss: 1.2171
- val_accuracy: 0.5310
Epoch 19/20
439/439 [=====] - 984s 2s/step - loss: 1.3661 - accuracy: 0.4809 - val_loss: 1.2117
- val_accuracy: 0.5325
Epoch 20/20
439/439 [=====] - 986s 2s/step - loss: 1.3398 - accuracy: 0.4886 - val_loss: 1.2429
- val_accuracy: 0.5152
```

Saving the trained model

```
In [37]: 1 model.save("model.h5")
2
```

Evaluate the model

```
In [38]: 1 # Evaluate the model
2 test_loss, test_acc = model.evaluate(test_generator)
3 print("validation accuracy :", str(test_acc*100)+"%")
4 print("validation loss :", test_loss)
```

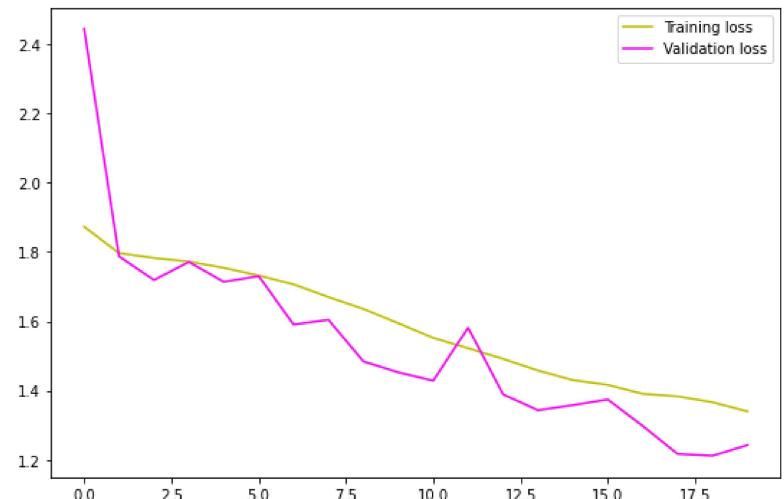
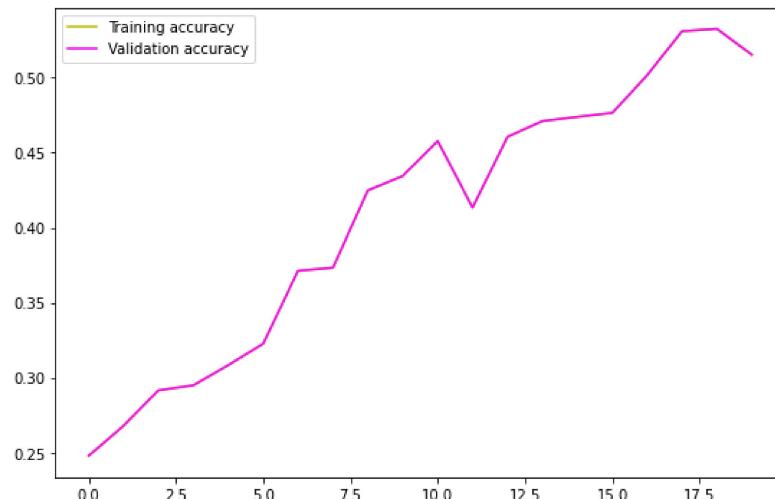
```
113/113 [=====] - 63s 561ms/step - loss: 1.2417 - accuracy: 0.5157
validation accuracy : 51.57425403594971%
validation loss : 1.241727352142334
```

Plotting Training and Validation plot

In [39]:

```
1 val_acc = history.history['val_accuracy']
2 loss = history.history['loss']
3 val_loss = history.history['val_loss']
4 epochs = range(len(val_acc))
5
6 fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(20, 6))
7 ax[0].plot(epochs, val_acc, 'y', label='Training accuracy')
8 ax[0].plot(epochs, val_acc, 'magenta', label='Validation accuracy')
9 ax[0].legend(loc=0)
10 ax[1].plot(epochs, loss, 'y', label='Training loss')
11 ax[1].plot(epochs, val_loss, 'magenta', label='Validation loss')
12 ax[1].legend(loc=0)
13
14 plt.suptitle('Training and validation')
15 plt.show()
```

Training and validation



Plotting the confusion matrix

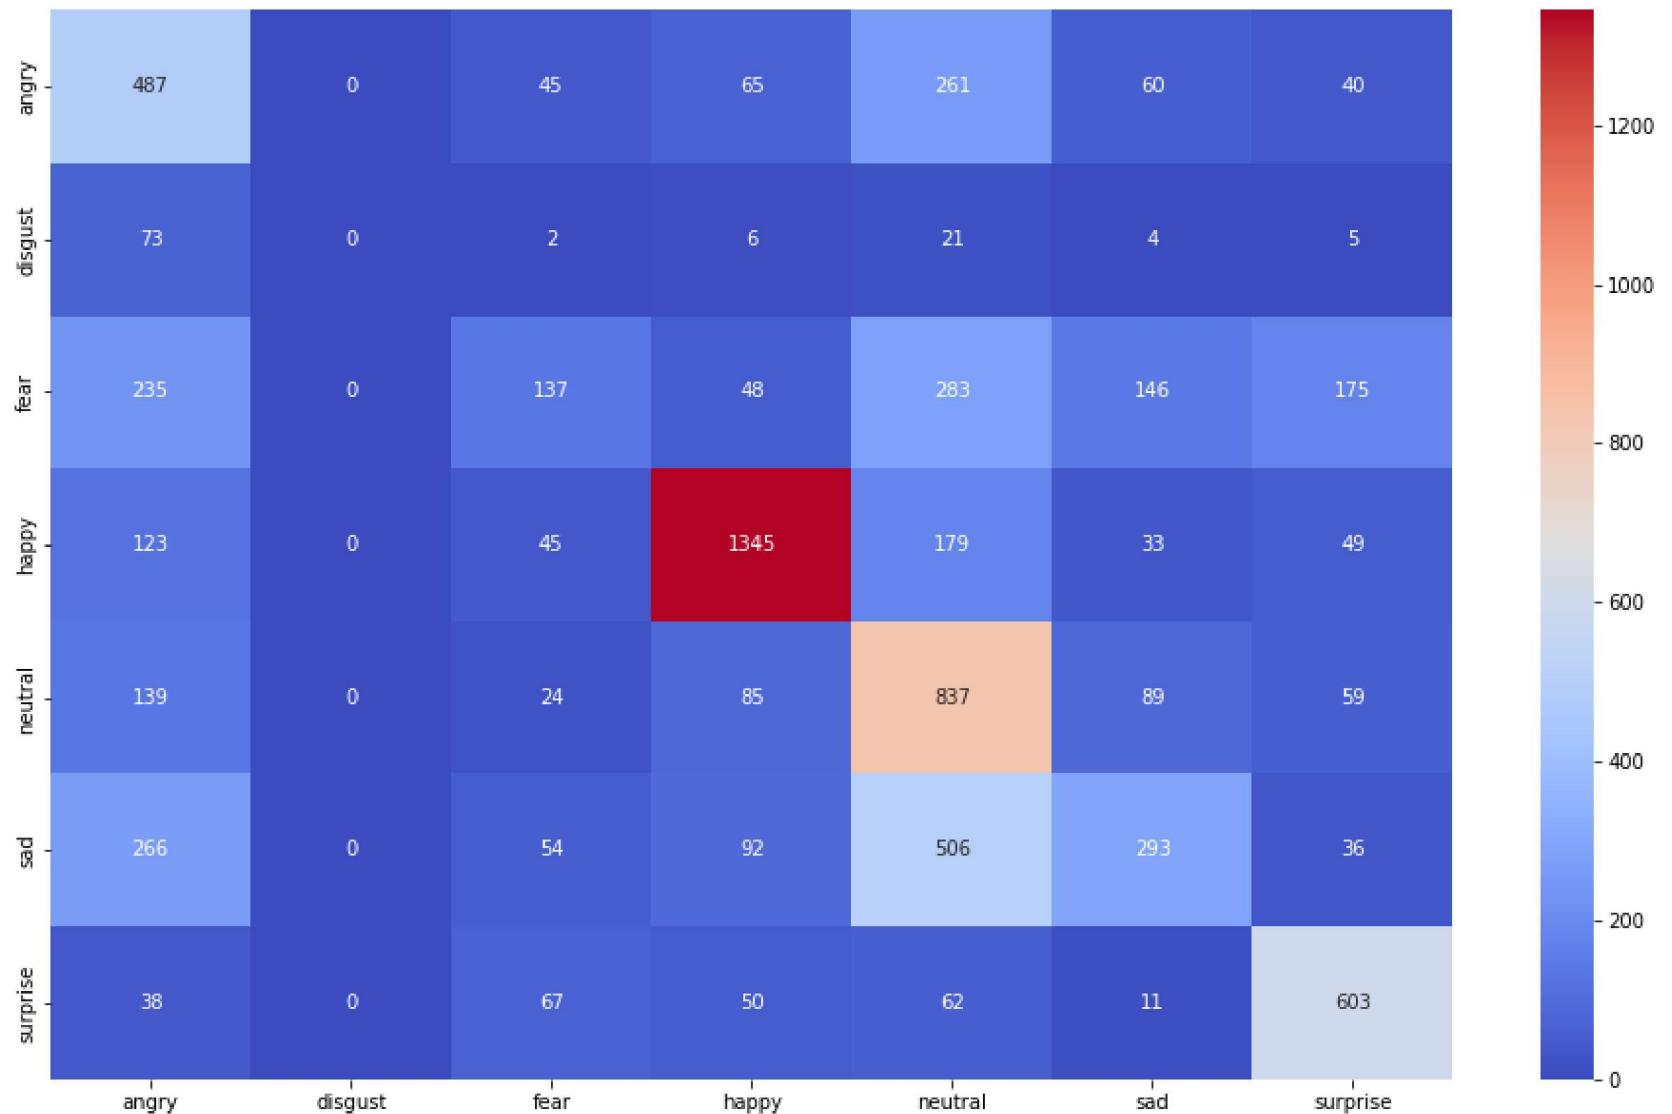
In [40]:

```
1 # confusion matrix
2 from sklearn.metrics import classification_report, confusion_matrix
3
4 y_pred = np.argmax(model.predict(test_generator), axis=-1)
5 print(classification_report(test_generator.classes, y_pred, target_names=test_generator.class_indices.keys()))
6
7 cm = confusion_matrix(test_generator.classes, y_pred)
8 plt.figure(figsize=(16,10))
9 sns.heatmap(cm, cmap=plt.cm.coolwarm, annot=True, fmt='.'0f', xticklabels=test_generator.class_indices.keys())
10 plt.show()
```

```
C:\Users\sbpat\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\metrics\_classification.py:
1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\sbpat\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\metrics\_classification.py:
1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\sbpat\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\metrics\_classification.py:
1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

angry	0.36	0.51	0.42	958
disgust	0.00	0.00	0.00	111
fear	0.37	0.13	0.20	1024
happy	0.80	0.76	0.78	1774
neutral	0.39	0.68	0.49	1233
sad	0.46	0.23	0.31	1247
surprise	0.62	0.73	0.67	831
accuracy			0.52	7178
macro avg	0.43	0.43	0.41	7178
weighted avg	0.52	0.52	0.49	7178



Testing our model with new image

In [41]:

```
1 # Testing our model with new image
2 image = cv2.imread("images/Sad.jpg")
3 from IPython.display import Image
4 Image(filename='Sad.jpg')
```

Out[41]:



In []:

```
1 # Model Prediction
2 import cv2
3 from tensorflow.python.keras.models import load_model
4 import os
5
6
7 # Load the trained model
8 model = load_model("model.h5")
9 # A list of emoticon categories
10 EMOTIONS = ['Angry', 'Disgust', 'Happy', 'Sad', 'Surprise', 'Neutral']
11 # Load image
12 img = image
13
14 # Trim the image to 48 x 48, and turn the grayscale image, normalization
15 frame = cv2.resize(img,(48,48),interpolation=cv2.INTER_BITS2)
16 gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY) / 255.0
17
18 # Reinvent the image dimension
19 gray = gray.reshape(1,48,48,1)
20
21 # Output the prediction
22 predicts = model.predict(gray)[0]
23 label = EMOTIONS[predicts.argmax()]
24 for (i,j) in zip(range(7),EMOTIONS):
25     predictss = predicts[i]
26     print("{:^10s}.".format(j)+"prediction rate is {0:.2f}%".format(predictss))
27 print( "\n\n The system considers this expression to be:",label)
28
```

Angry prediction rate is 0.24%

Disgust prediction rate is 0.00%

Happy prediction rate is 0.08%

Sad prediction rate is 0.27%

Surprise prediction rate is 0.26%

Neutral prediction rate is 0.10%

The system considers this expression to be: Sad

Song Recommdation

In []:

```
1 # Song Recommdation
2 if (label=='Angry'):
3     path="test\\Angry\\"
4     files=os.listdir(path)
5     d=random.choice(files)
6     print("Now Playing:",d)
7     audio = Audio(filename='song\\Angry\\'+ d,autoplay=True)
8     display(audio)
9
10 elif (label=='Disgust'):
11     path="test\\Disgust\\"
12     files=os.listdir(path)
13     d=random.choice(files)
14     print("Now Playing:",d)
15     audio = Audio(filename='song\\Disgust\\'+ d,autoplay=True)
16     display(audio)
17
18 elif (label=="Happy"):
19     path="test\\happy\\"
20     files=os.listdir(path)
21     d=random.choice(files)
22     print("Now Playing:",d)
23     audio = Audio(filename='song\\Happy\\'+ d,autoplay=True)
24     display(audio)
25
26 elif (label=='Sad'):
27     path="test//sad//"
28     files=os.listdir(path)
29     d=random.choice(files)
30     print("Now Playing:",d)
31     audio = Audio(filename='song\\Sad\\'+ d,autoplay=True)
32     display(audio)
33 elif (label=='Surprise'):
34     path="test\\Surprise\\"
35     files=os.listdir(path)
36     d=random.choice(files)
37     print("Now Playing:",d)
38     audio = Audio(filename='song\\Surprise\\'+ d,autoplay=True)
39     display(audio)
40
41 elif (label=='Neutral'):
42     path="test\\Neutral\\"
43     files=os.listdir(path)
```

```
44     d=random.choice(files)
45     print("Now Playing:",d)
46     audio = Audio(filename='song\\Neutral\\'+ d,autoplay=True)
47     display(audio)
48
```

Now Playing: English Sad Song Slow Music heart touching painful.mp3

Thank You !

In []:

1